Training split, recorded labels only, before any predictive model is trained
1 · Diagnose
Causal discovery
+ assumption gate
(POINT / BOUNDED /
REFUSED)
2 · Propose
Fixed rule: protected,
labeling-process
variables, their neighbors,
the label's neighbors
3 · Check
Recorded vs expected
outcomes (O/E), only for
groups tied to the
labeling process,
not to the outcome
4 · Flag
Groups with fewer
recorded outcomes
than their risk
factors predict
Auditor
Verify labels in
the flagged groups